


Linear Haskell for string builders

Andrew Lelechenko
`andrew.lelechenko@gmail.com`

MuniHac, 13.09.2025
<https://tinyurl.com/bdh8wyyp>





List concatenation is linear ...

```
type String = [Char]
```

```
output :: String
```

```
output = longString ++ veryLongString ++ extraLongString
```

```
output :: String
```

```
output = (longString ++ veryLongString) ++ extraLongString
```

```
infixr 5  ++
```

```
(++) :: [a] → [a] → [a]
```

```
(++) []      ys = ys
```

```
(++) (x : xs) ys = x : xs ++ ys
```

List concatenation is linear only in its first argument

```
type String = [Char]
```

```
good :: String
```

```
good = longString ++ veryLongString ++ extraLongString
```

```
good' :: String
```

```
good' = longString ++ (veryLongString ++ extraLongString)
```

```
bad :: String
```

```
bad = (longString ++ veryLongString) ++ extraLongString
```

```
ghci> length (replicate 100000000 'x' ++ ("foo" ++ "bar"))  
100000006
```

```
(1.18 secs, 11,200,068,064 bytes)
```

```
ghci> length ((replicate 100000000 'x' ++ "foo") ++ "bar")  
100000006
```

```
(1.84 secs, 16,800,068,064 bytes)
```

How to define instances of class Show?

```
data User = User
  { name      :: String
  , address  :: String
  , phone     :: String }
```

```
instance Show User where
  show User{..} = name ++ address ++ phone
```

```
data Order = Order
  { user       :: User
  , product   :: String
  , date      :: String }
```

```
instance Show Order where
  show Order{..} = show user ++ product ++ date
```

Define showsPrec, not show

```
class Show a where
  show :: a → String
  showsPrec :: Int → a → (String → String)
  {-# MINIMAL show | showsPrec #-}

instance Show User where
  showsPrec _ User{..} rest =
    name ++ address ++ phone ++ rest

instance Show Order where
  showsPrec p Order{..} rest =
    showsPrec p user (product ++ date ++ rest)
```

Compose functions instead of concatenating data

```
instance Show User where
  -- showsPrec _ User{..} rest =
  --   name ++ address ++ phone ++ rest
  showsPrec _ User{..} =
    (name ++) . (address ++) . (phone ++)

instance Show Order where
  -- showsPrec p Order{..} rest =
  --   showsPrec p user (product ++ date ++ rest)
  showsPrec p Order{..} =
    showsPrec p user . (product ++) . (date ++)
```

Which is faster?

```
(long ++) . ((veryLong ++) . (extraLong ++)) $ []
```

```
((long ++) . (veryLong ++)) . (extraLong ++) $ []
```

```
((long ++) . (veryLong ++)) . (extraLong ++) $ []  
  = ((long ++) . (veryLong ++)) $ (extraLong ++) $ []  
    = (long ++) $ (veryLong ++) $ (extraLong ++) $ []
```


Builders for lists

Efficient concatenation is possible, but requires diligence to add new chunks from the left only.

```
newtype Builder = Builder (String → String)
```

```
fromString :: String → Builder  
fromString xs = Builder (xs ++)
```

```
toString :: Builder → String  
toString (Builder f) = f []
```

```
instance Semigroup Builder where  
    Builder f <> Builder g = Builder (f . g)
```

Builder allows to concatenate left and right, although has an increased constant factor.

~~String~~

StrictText

Concatenation of StrictText is linear ...

```
data StrictText = Text
  { buffer :: ByteArray
  , offset :: Int
  , length :: Int }
  -- len(buffer) could be /= offset + length

concatRight :: StrictText
concatRight = longText ++ (veryLongText ++ extraLongText)

concatLeft :: StrictText
concatLeft = (longText ++ veryLongText) ++ extraLongText
```

Concatenation of StrictText is linear in both arguments

```
data StrictText = Text
  { buffer :: ByteArray
  , offset :: Int
  , length :: Int }
  -- len(buffer) could be /= offset + length

concatRight :: StrictText
concatRight = longText ++ (veryLongText ++ extraLongText)

concatLeft :: StrictText
concatLeft = (longText ++ veryLongText) ++ extraLongText
```

Data.Text.Lazy.Builder sidesteps the issue

```
data LazyText = Empty | Chunk StrictText LazyText
```

```
newtype Builder = Builder {  
  ∀ s. (Buffer s → ST s [StrictText])  
    → (Buffer s → ST s [StrictText]) }  
}
```

```
data Buffer s = Buffer  
  { buffer :: MutableByteArray s  
  , offset :: Int  
  , used    :: Int  
  , unused  :: Int }  
-- len(buffer) = offset + used + unused
```

```
data TextBuilder = TextBuilder
  -- Estimated maximum size of the byte array to allocate.
  Int
  -- Function that populates a preallocated bytearray
  -- of the estimated maximum size specified above provided
  -- an offset into it and producing the offset after.
  (∀ s. MutableByteArray s → Int → ST s Int)

instance Semigroup TextBuilder where
  TextBuilder lenL writeL <> TextBuilder lenR writeR =
    TextBuilder
      (lenL + lenR)
      (\array offset → do
        offsetAfter1 ← writeL array offset
        writeR array offsetAfter1
      )
```

Java-style string builder

```
data Buffer = Buffer
  { buffer :: ByteArray
  , used   :: Int }
```

```
(++) :: Buffer → StrictText → Buffer
```

```
Buffer arr used ++ Text srcArr srcOff srcLen = runST $ do
  let unused = sizeofByteArray arr - used
  if unused ≥ srcLen then do
    mutArr ← unsafeThawByteArray arr
    copyByteArray mutArr used srcArr srcOff srcLen
    arr' ← unsafeFreezeByteArray mutArr
    pure $ Buffer arr' (used + srcLen)
  else do
    mutArr ← newByteArray ((used + srcLen) * 2)
    copyByteArray mutArr 0 arr 0 used
    copyByteArray mutArr used srcArr srcOff srcLen
    arr' ← unsafeFreezeByteArray mutArr
```

Honest mutable Buffer

```
data MutBuffer s = MutBuffer
  { buffer :: MutableByteArray s
  , used   :: Int }

(++ ) :: MutBuffer s → StrictText → ST s (MutBuffer s)
MutBuffer mutArr used ++ Text srcArr srcOff srcLen = do
  size ← getSizeOfMutableByteArray mutArr
  let unused = size - used
  if unused ≥ srcLen then do
    copyByteArray mutArr used srcArr srcOff srcLen
    pure $ MutBuffer mutArr (used + srcLen)
  else do
    let newSize = (used + srcLen) * 2
    mutArr' ← resizeMutableByteArray mutArr newSize
    copyByteArray mutArr' used srcArr srcOff srcLen
    pure $ MutBuffer mutArr' (used + srcLen)
```


attoparsec used linear types before linear types

```
data Builder = Builder
  { gen      :: Int
  , buffer  :: ByteArray -- ^ also stores 'gen' at start
  , used    :: Int }
```

Commit 62856d6 by @bos on May 30, 2014

The fact of having a mutable buffer really helps with performance, but ... it does have a consequence: if someone misuses [it] ... they could overwrite data.

... we use two generation counters (one mutable, one immutable) to track the number of appends to a mutable buffer. If the counters ever get out of sync, someone is appending twice to a mutable buffer, so we duplicate the entire buffer in order to preserve the immutability of its older self.

While we could go a step further and gain protection against API abuse on a multicore system, by use of an atomic increment instruction to bump the mutable generation counter, that would be very expensive. ... Clients should never call a continuation more than once; **we lack a linear type system** that could enforce this. ...

Linear buffer API — 1

```
data Buffer :: TYPE ('BoxedRep 'Unlifted) where
  Buffer :: {-# UNPACK #-} !Text → Buffer
```

appendBounded

```
:: Int -- ^ Upper bound for the number of bytes to write
→ (∀ s. MutableByteArray s → Int → ST s Int)
-- ^ Write bytes starting from the given offset
-- and return an actual number of bytes written.
→ Buffer → Buffer
```

appendExact

```
:: Int -- ^ Exact number of bytes to write
→ (∀ s. A.MArray s → Int → ST s ())
-- ^ Write bytes starting from the given offset
→ Buffer → Buffer
```

```
(|>) :: Buffer → StrictText → Buffer
```

Linear buffer API — 2

$(|>) :: \text{Buffer} \multimap \text{StrictText} \rightarrow \text{Buffer}$

$(<|) :: \text{StrictText} \rightarrow \text{Buffer} \multimap \text{Buffer}$

$(><) :: \text{Buffer} \multimap \text{Buffer} \multimap \text{Buffer}$

```
> runBuffer (\b → b |> "foo" |> "bar")  
"foobar"
```

$(|>\#) :: \text{Buffer} \multimap \text{Addr\#} \rightarrow \text{Buffer}$

$(\#<|) :: \text{Addr\#} \rightarrow \text{Buffer} \multimap \text{Buffer}$

```
> runBuffer (\b → b |>\# "foo"\# |>\# "bar"\#)  
"foobar"
```

$(|>.) :: \text{Buffer} \multimap \text{Char} \rightarrow \text{Buffer}$

$(.<|) :: \text{Char} \rightarrow \text{Buffer} \multimap \text{Buffer}$

```
> runBuffer (\b → b |>. 'q' |>. 'w')  
"qw"
```

Linear buffer API — 3

```
(|>$) :: (Integral a, FiniteBits a) => Buffer -> a -> Buffer
($<|) :: (Integral a, FiniteBits a) => a -> Buffer -> Buffer
> runBuffer (\b -> b |>$ (42 :: Int))
"42"
```

[illegible]

```
(|>&) :: (Integral a, FiniteBits a) => Buffer -> a -> Buffer
(&<|) :: (Integral a, FiniteBits a) => a -> Buffer -> Buffer
> runBuffer (\b -> b |>& (42 :: Int))
"2a"
```

```
(|>%) :: Buffer -> Double -> Buffer
(%<|) :: Double -> Buffer -> Buffer
> runBuffer (\b -> b |>% 123.456)
"123.456"
```

No linear types? No problem

```
newtype Builder = Builder { unBuilder :: Buffer  $\multimap$  Buffer }
```

```
fromText :: Text  $\rightarrow$  Builder
```

```
fromChar :: Char  $\rightarrow$  Builder
```

```
fromAddr :: Addr#  $\rightarrow$  Builder
```

```
fromDec :: (Integral a, FiniteBits a)  $\Rightarrow$  a  $\rightarrow$  Builder
```

```
fromUnboundedDec :: Integral a  $\Rightarrow$  a  $\rightarrow$  Builder
```

```
fromHex :: (Integral a, FiniteBits a)  $\Rightarrow$  a  $\rightarrow$  Builder
```

```
fromDouble :: Double  $\rightarrow$  Builder
```

```
runBuilder :: Builder  $\rightarrow$  StrictText
```

```
runBuilderBS :: Builder  $\rightarrow$  StrictByteString
```

Benchmarks with GHC 9.12 on aarch64

	text	text-builder		This package	
Text					
1000	80.5 μ s	26.7 μ s	0.33x	23.1 μ s	0.29x
1000000	216 ms	107 ms	0.49x	22.9 ms	0.11x
Char					
1000	35.4 μ s	18.4 μ s	0.52x	7.68 μ s	0.22x
1000000	175 ms	178 ms	1.02x	10.5 ms	0.06x
Decimal					
1000	148 μ s	738 μ s	5.00x	106 μ s	0.72x
1000000	334 ms	2.803 s	8.40x	108 ms	0.32x
Hexadecimal					
1000	862 μ s	141 μ s	0.16x	44.6 μ s	0.05x
1000000	1.502 s	228 ms	0.15x	45.9 ms	0.03x
Double					
1000	14.2 ms	71.9 ms	5.05x	671 μ s	0.05x
1000000	14.366 s	101.342 s	7.05x	689 ms	0.05x

Thank you!

 Bodigrim

@ andrew.lelechenko@gmail.com

 github.com/Bodigrim/linear-builder

» hackage.haskell.org/package/text-builder-linear