

Linear Haskell for string builders

Andrew Lelechenko
`andrew.lelechenko@gmail.com`

MuniHac, 13.09.2025

List concatenation is linear ...

```
type String = [Char]
```

```
concatRight :: String
```

```
concatRight = long ++ veryLong ++ extraLong
```

```
concatLeft :: String
```

```
concatLeft = (long ++ veryLong) ++ extraLong
```

```
infixr 5  ++
```

```
(++) :: [a] → [a] → [a]
```

```
(++) []      ys = ys
```

```
(++) (x : xs) ys = x : (xs ++ ys)
```

List concatenation is linear only in its first argument

```
> length (replicate 100000000 'x' ++ ("foo" ++ "bar"))
100000006
(1.18 secs, 11,200,068,064 bytes)
> length ((replicate 100000000 'x' ++ "foo") ++ "bar")
100000006
(1.84 secs, 16,800,068,064 bytes)

> length (foldr (++) [] (map (:[ ]) ['\0'..'\10000']))
10001
(0.01 secs, 2,945,344 bytes)
> length (foldl (++) [] (map (:[ ]) ['\0'..'\10000']))
10001
(0.84 secs, 4,299,915,184 bytes)
```

How to define instances of class Show?

```
data User = User
  { name      :: String
  , address  :: String
  , phone    :: String }
```

```
instance Show User where
  show User{..} = name ++ address ++ phone
```

```
data Order = Order
  { user       :: User
  , product   :: String
  , date      :: String }
```

```
instance Show Order where
  show Order{..} = show user ++ product ++ date
```

Define showsPrec, not show

```
class Show a where
  show :: a → String
  showsPrec :: Int → a → (String → String)
  {-# MINIMAL show | showsPrec #-}

instance Show User where
  showsPrec _ User{..} rest =
    name ++ address ++ phone ++ rest

instance Show Order where
  showsPrec p Order{..} rest =
    showsPrec p user (product ++ date ++ rest)
```

Compose functions instead of concatenating data

```
instance Show User where
  -- showsPrec _ User{..} rest =
  --   name ++ address ++ phone ++ rest
  showsPrec _ User{..} =
    (name ++) . (address ++) . (phone ++)

instance Show Order where
  -- showsPrec p Order{..} rest =
  --   showsPrec p user (product ++ date ++ rest)
  showsPrec p Order{..} =
    showsPrec p user . (product ++) . (date ++)
```

Which is faster?

```
(long ++) . ((veryLong ++) . (extraLong ++)) $ []
```

vs.

```
((long ++) . (veryLong ++)) . (extraLong ++) $ []
```

There is no big difference!

```
((long ++) . (veryLong ++)) . (extraLong ++) $ []  
  = ((long ++) . (veryLong ++)) $ (extraLong ++) $ []  
    = (long ++) $ (veryLong ++) $ (extraLong ++) $ []
```

List concatenation: recap

Efficient concatenation is possible, but requires diligence to add new chunks from the left side only.

```
newtype DList = DList (String → String)
```

```
fromString :: String → DList  
fromString xs = DList (xs ++)
```

```
toString :: DList → String  
toString (DList f) = f []
```

```
instance Semigroup DList where  
    DList f <> DList g = DList (f . g)
```

DList allows to concatenate left and right, although has an increased constant factor.

~~String~~ (Strict)Text

Concatenation of Text is linear ...

```
data Text = Text
  { buffer :: ByteArray
  , offset :: Int
  , length :: Int }
  -- len(buffer) could be /= offset + length

concatRight :: Text
concatRight = longText <> (veryLongText <> extraLongText)

concatLeft :: Text
concatLeft = (longText <> veryLongText) <> extraLongText
```

Concatenation of Text is linear in both arguments

```
> let x = T.replicate 100000000 "x" in T.length (x <> (x <> x))
300000000
(0.14 secs, 600,067,864 bytes)
> let x = T.replicate 100000000 "x" in T.length ((x <> x) <> x)
300000000
(0.12 secs, 600,067,864 bytes)

> T.length (foldr (<>) mempty (map T.singleton ['\0'..'\10000']))
10001
(0.16 secs, 151,491,368 bytes)
> T.length (foldl (<>) mempty (map T.singleton ['\0'..'\10000']))
10001
(0.14 secs, 133,854,104 bytes)
```

Efficient concatenation requires us to guess the size of the final result

Data.Text.Lazy.Builder sidesteps the issue

```
newtype Builder = Builder {  
  ∀ s. (Buffer s → ST s [Text])  
    → (Buffer s → ST s [Text]) }  
  
data Buffer s = Buffer  
  { buffer :: MutableByteArray s  
  , offset :: Int  
  , used    :: Int  
  , unused  :: Int }  
-- len(buffer) = offset + used + unused
```

```
data TextBuilder = TextBuilder
  -- Estimated max size of the bytearray to allocate.
  Int
  -- Function that populates a preallocated bytearray
  -- of the estimated max size specified above provided
  -- an offset into it and producing the offset after.
  (∀ s. MutableByteArray s → Int → ST s Int)
```

```
instance Semigroup TextBuilder where
  TextBuilder lenL writeL <> TextBuilder lenR writeR =
    TextBuilder
      (lenL + lenR)
      (\array offset → do
        offsetAfter1 ← writeL array offset
        writeR array offsetAfter1
      )
```

Java-style string builder

```
data Buffer = Buffer
  { buffer :: ByteArray, used :: Int }

(++ ) :: Buffer → Text → Buffer
Buffer arr used ++ Text srcArr srcOff srcLen = runST $ do
  let unused = sizeofByteArray arr - used
  if unused ≥ srcLen then do
    mutArr ← unsafeThawByteArray arr
    copyByteArray mutArr used srcArr srcOff srcLen
    arr' ← unsafeFreezeByteArray mutArr
    pure $ Buffer arr' (used + srcLen)
  else do
    mutArr ← newByteArray ((used + srcLen) * 2)
    copyByteArray mutArr 0 arr 0 used
    copyByteArray mutArr used srcArr srcOff srcLen
    arr' ← unsafeFreezeByteArray mutArr
    pure $ Buffer arr' (used + srcLen)
```

Honest mutable Buffer

```
data MutBuffer s = MutBuffer
  { buffer :: MutableByteArray s, used :: Int }

(++ ) :: MutBuffer s → Text → ST s (MutBuffer s)
MutBuffer mutArr used ++ Text srcArr srcOff srcLen = do
  size ← getSizeofMutableByteArray mutArr
  let unused = size - used
  if unused ≥ srcLen then do
    copyByteArray mutArr used srcArr srcOff srcLen
    pure $ MutBuffer mutArr (used + srcLen)
  else do
    let newSize = (used + srcLen) * 2
    mutArr' ← resizeMutableByteArray mutArr newSize
    copyByteArray mutArr' used srcArr srcOff srcLen
    pure $ MutBuffer mutArr' (used + srcLen)
```

attoparsec used linear types before linear types

```
data Builder = Builder
  { gen      :: Int
  , buffer  :: ByteArray -- ^ also stores 'gen' at start
  , used    :: Int }
```

Commit 62856d6 by @bos on May 30, 2014

The fact of having a mutable buffer really helps with performance, but ... it does have a consequence: if someone misuses [it] ... they could overwrite data.

... we use two generation counters (one mutable, one immutable) to track the number of appends to a mutable buffer. If the counters ever get out of sync, someone is appending twice to a mutable buffer, so we duplicate the entire buffer in order to preserve the immutability of its older self.

While we could go a step further and gain protection against API abuse on a multicore system, by use of an atomic increment instruction to bump the mutable generation counter, that would be very expensive. ... Clients should never call a continuation more than once; **we lack a linear type system** that could enforce this. ...

Linear and unlifted types for the rescue

```
data Buffer :: TYPE ('BoxedRep 'Unlifted) where
  Buffer :: {-# UNPACK #-} !Text → Buffer
  -- ^ constructor is not exported
```

appendBounded

```
:: Int -- ^ Upper bound for the number of bytes to write
→ (∀ s. MutableByteArray s → Int → ST s Int)
-- ^ Write bytes starting from the given offset
-- and return an actual number of bytes written.
→ Buffer → Buffer
```

```
(▷) :: Buffer → Text → Buffer
```

```
runBuffer :: (Buffer → Buffer) → Text
```

```
runBufferBS :: (Buffer → Buffer) → StrictByteString
```

Appending letters

$(\triangleright) :: \text{Buffer} \multimap \text{Text} \rightarrow \text{Buffer}$

$(\triangleleft) :: \text{Text} \rightarrow \text{Buffer} \multimap \text{Buffer}$

$(><) :: \text{Buffer} \multimap \text{Buffer} \multimap \text{Buffer}$

```
> runBuffer (\b → b ▷ "foo" ▷ "bar")  
"foobar"
```

$(\triangleright\#) :: \text{Buffer} \multimap \text{Addr}\# \rightarrow \text{Buffer}$

$(\#\triangleleft) :: \text{Addr}\# \rightarrow \text{Buffer} \multimap \text{Buffer}$

```
> runBuffer (\b → b ▷# "foo"# ▷# "bar"#)  
"foobar"
```

$(\triangleright.) :: \text{Buffer} \multimap \text{Char} \rightarrow \text{Buffer}$

$(. \triangleleft) :: \text{Char} \rightarrow \text{Buffer} \multimap \text{Buffer}$

```
> runBuffer (\b → b ▷. 'q' ▷. 'w')  
"qw"
```

Appending numbers

```
(▷$) :: (Integral a, FiniteBits a) ⇒ Buffer → a → Buffer
($◁) :: (Integral a, FiniteBits a) ⇒ a → Buffer → Buffer
> runBuffer (\b → b ▷$ (42 :: Int))
"42"
```

[illegible]

```
(▷&) :: (Integral a, FiniteBits a) ⇒ Buffer → a → Buffer
(&◁) :: (Integral a, FiniteBits a) ⇒ a → Buffer → Buffer
> runBuffer (\b → b ▷& (42 :: Int))
"2a"
```

```
(▷%) :: Buffer → Double → Buffer
(%)< :: Double → Buffer → Buffer
> runBuffer (\b → b ▷% 123.456)
"123.456"
```

No linear types? No problem

```
newtype Builder = Builder { unBuilder :: Buffer  $\multimap$  Buffer }
```

```
fromText :: Text  $\rightarrow$  Builder
```

```
fromChar :: Char  $\rightarrow$  Builder
```

```
fromAddr :: Addr#  $\rightarrow$  Builder
```

```
fromDec :: (Integral a, FiniteBits a)  $\Rightarrow$  a  $\rightarrow$  Builder
```

```
fromUnboundedDec :: Integral a  $\Rightarrow$  a  $\rightarrow$  Builder
```

```
fromHex :: (Integral a, FiniteBits a)  $\Rightarrow$  a  $\rightarrow$  Builder
```

```
fromDouble :: Double  $\rightarrow$  Builder
```

```
runBuilder :: Builder  $\rightarrow$  Text
```

```
runBuilderBS :: Builder  $\rightarrow$  StrictByteString
```

Benchmarks with GHC 9.12 on aarch64

	text	text-builder		This package	
Text					
1000	80.5 μ s	26.7 μ s	0.33x	23.1 μ s	0.29x
1000000	216 ms	107 ms	0.49x	22.9 ms	0.11x
Char					
1000	35.4 μ s	18.4 μ s	0.52x	7.68 μ s	0.22x
1000000	175 ms	178 ms	1.02x	10.5 ms	0.06x
Decimal					
1000	148 μ s	738 μ s	5.00x	106 μ s	0.72x
1000000	334 ms	2.803 s	8.40x	108 ms	0.32x
Hexadecimal					
1000	862 μ s	141 μ s	0.16x	44.6 μ s	0.05x
1000000	1.502 s	228 ms	0.15x	45.9 ms	0.03x
Double					
1000	14.2 ms	71.9 ms	5.05x	671 μ s	0.05x
1000000	14.366 s	101.342 s	7.05x	689 ms	0.05x

Thank you!

 Bodigrim

@ andrew.lelechenko@gmail.com

 github.com/Bodigrim/linear-builder

» hackage.haskell.org/package/text-builder-linear